



NIGHTMARE

Building autonomous security research that is auditable by design.

PHANTOMBUGZ • SEED ROUND • 2026

THE PROBLEM

Software ships faster than it can be secured.

- ◆ Vulnerability research is manual, slow, and expensive.
- ◆ Skilled offensive-security talent is scarce and does not scale.
- ◆ Attackers already automate. Most defenders still hunt by hand.
- ◆ Every audit that never happens is an open door.


THE SOLUTION

Meet Nightmare.

An evidence-first platform under development for authorized, autonomous security research.

- ◆ Orchestrates specialized AI analysis over code and controlled targets.
- ◆ Routes candidate tests through policy-gated sandbox execution.
- ◆ Captures evidence, decisions, and provenance in a replayable audit trail.
- ◆ Targets reproducible findings and repair guidance with less manual triage.

A governed loop from authorized target to evidence-backed report.



The product goal is not another suggestion engine. It is a governed system that preserves the evidence needed to verify every result.

BUILT TODAY

The verified foundation is in place.

Nightmare Slice Zero is merged with a recorded 15-of-15 acceptance result. It establishes the control plane required before model-driven security testing is added.

- ◆ Signed structured inputs with bounded parsing and explicit lifecycle controls.
- ◆ Authority separation across proposal, evidence ingestion, execution, and governance.
- ◆ Atomic hash-chained records plus a replayable graph projection for audit.

TRACTION

Real recurring revenue, growing fast.

\$100k-\$499k

Annual Recurring Revenue

100-149%

Year-over-year growth

Seed

Current fundraising stage

Company-reported figures supplied for this raise. Revenue validates the broader PhantomBugz security business while Nightmare advances from verified foundation to product.

Start with buyers who already pay for deeper security validation.

Researchers

Bug bounty and independent security teams seeking higher throughput

Pentest

Firms that need repeatable evidence and more analyst leverage

Enterprise

Application security teams moving toward continuous validation

The beachhead is authorized offensive-security work where reproducibility, auditability, and researcher time have clear economic value.

BUSINESS MODEL

A planned recurring model that scales with usage.

HOW WE CHARGE

- ◆ Per-seat subscriptions for security teams and researchers.
- ◆ Platform licenses for enterprise deployment.
- ◆ Usage-based compute for large autonomous runs.

EXPANSION PATH

- ◆ Begin with controlled research and evaluation workflows.
- ◆ Expand usage as supported targets and evidence quality improve.
- ◆ Grow within accounts as teams adopt continuous validation.

WHY NOW

The capability, the compute, and the need arrived at the same time.

- ◆ Modern LLMs can finally reason over real code at long context.
- ◆ NVIDIA GPU infrastructure makes long-context inference and evaluation practical.
- ◆ Software and attack surface are exploding while talent stays flat.
- ◆ Evidence-first architecture can turn probabilistic output into an auditable workflow.

Auditability is designed into the control plane.

- ◆ Signed intermediate representation and bounded parsing for well-formed tasks.
- ◆ Separated authority for proposals, evidence, execution, and governance.
- ◆ Atomic hash-chained ledger behavior with deterministic replay.
- ◆ Architecture designed to produce traceable evaluation and training data as model specialists come online.

DIFFERENTIATION

Move from untracked suggestions to governed evidence.

THE ALTERNATIVES

- ◆ Legacy scanners often create large manual triage queues.
- ◆ Manual pentests deliver depth but have limited throughput.
- ◆ Generic AI copilots can generate suggestions without a durable evidence trail.

NIGHTMARE DESIGN

- ◆ Policy-gated orchestration from scoped task to report.
- ◆ Evidence and provenance captured as first-class artifacts.
- ◆ Replayable architecture designed for NVIDIA GPU scale.

Founder-led today, specialist expansion after the raise.

PhantomBugz is building from hands-on security research and a verified systems foundation, with a focused seed hiring plan.

- ◆ Founder-led product direction and security-research execution.
- ◆ First technical hires: ML systems and security research engineering.
- ◆ Commercial priority: design partners and enterprise product readiness.

Prove in controlled workflows, then expand.

- ◆ Recruit design partners among researchers and boutique pentest firms.
- ◆ Benchmark end-to-end reproducibility and analyst time saved.
- ◆ Package the workflow for enterprise application-security teams.
- ◆ Use successful pilots as references for repeatable sales.

THE ASK

Raising a **\$250k to \$1M** seed round.

Compute

Benchmark and deploy model specialists on NVIDIA GPU infrastructure

Team

Key engineering and security-research hires

Pilots

Design-partner evaluation across research, pentest, and enterprise

Target milestones: an end-to-end sandbox benchmark, measured reproducibility and evidence quality, and a pilot-ready product.



**Make offensive security autonomous,
continuous, and verifiable.**

PhantomBugz • simon@phantombugz.com • phantombugz.com